

# Homework 3: Inference & Learning for Temporal State Space Models

Brown University CS242: Probabilistic Graphical Models, Fall 2014

## Official Solutions

### Question 1:

a) *Figure 1 shows the Kalman filter estimates with two standard deviations of error.*


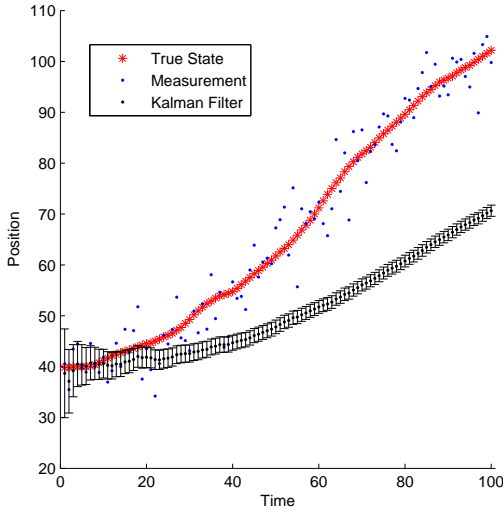
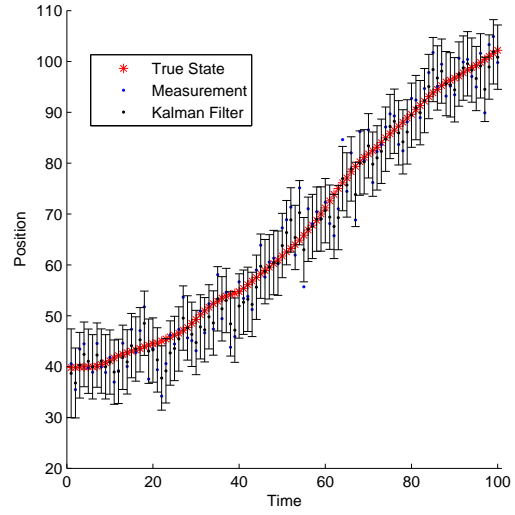


Figure 1: Kalman filter output for constant velocity model.

- b) *Figure 2 shows the Kalman filter estimates with two standard deviations of error for  $\sigma_x^2 = 0.01/3$ , and  $\sigma_x^2 = 10$ . The primary difference between this and Figure 1 is that when using the correct dynamics noise ( $\sigma_x^2 = 0.01/3$ ), but incorrectly assuming that the velocity is zero, the filter produces inaccurate state estimates. Using a dynamics model that allows for bigger movements ( $\sigma_x^2 = 10$ ), the estimates produced by the Kalman filter follow roughly the true state, but the standard deviation on the estimates is larger. Intuitively, the overly simplified constant-position model must treat the unmodeled velocity as noise.*
- c) *See solution code in `particle_filter.m`.*
- d) *Figure 3(a) shows the particle filter compared to the Kalman filter for three random initializations with 100 particles. Particle filter state estimates are generally close to the true posterior means.*



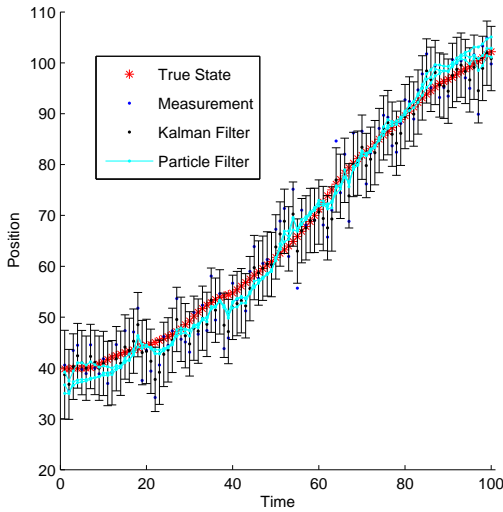
(a)




(b)

Figure 2: (a) Kalman filter output for constant position model with  $\sigma_x^2 = 0.01/3$ . (b) Kalman filter output for constant position model with  $\sigma_x^2 = 10$

e) Figure 3(b) shows the particle filter compared to the Kalman filter for three random instances with 20 particles. In many cases, the state estimates deviate significantly from the true posterior means. The particle filter is only robust if sufficiently many particles are used; for this model and data, 20 appears too small.



(a)



(b)

Figure 3: Particle filter output for three random initializations using (a) 100 particles, and (b) 20 particles.




Figure 4: Comparison of Kalman Filter and Particle Filter with 1000 particles on constant velocity model with outliers. The observation likelihood of the Kalman filter does not account for outliers, whereas the particle filter is matched to the model.

f) Figure 4 shows the results of tracking on the constant velocity model with outliers. We see that the particle filter, with matched observation model, is more robust to the noisy data.

### Question 2:

- a) Figure 5 shows a comparison between the Kalman Filter and Smoother. In general the smoother reduces covariance and produces a better estimate of the latent state.
- b) We can introduce, and integrate over, the latent state  $x_{t+1}$  as,

$$p(y_{t+1} | y_1, \dots, y_t) = \int_{\mathcal{X}_{t+1}} \int_{\mathcal{X}_t} p(y_{t+1} | x_{t+1}) p(x_{t+1} | x_t) p(x_t | y_1, \dots, y_t) dx_t dx_{t+1}$$

Observe that  $p(y_{t+1} | x_{t+1}) = N(y_{t+1} | Cx_{t+1}, R)$  is given by the observation model,  $p(x_{t+1} | x_t) = N(x_{t+1} | Ax_t, Q)$  the dynamics, and  $p(x_t | y_1, \dots, y_t) = N(x_t | \hat{x}_{t|t}, P_{t|t})$  is computed by the Kalman filter. Using standard Gaussian identities we can solve,

$$\begin{aligned} p(y_{t+1} | y_1, \dots, y_t) &= \int_{\mathcal{X}_{t+1}} \int_{\mathcal{X}_t} N(y_{t+1} | Cx_{t+1}, R) N(x_{t+1} | Ax_t, Q) N(x_t | \hat{x}_{t|t}, P_{t|t}) dx_t dx_{t+1} \\ &= \int_{\mathcal{X}_{t+1}} N(y_{t+1} | Cx_{t+1}, R) N(x_{t+1} | A\hat{x}_{t|t}, Q + AP_{t|t}A^T) dx_{t+1} \\ &= N(y_{t+1} | CA\hat{x}_{t|t}, R + C(Q + AP_{t|t}A^T)C^T). \end{aligned}$$

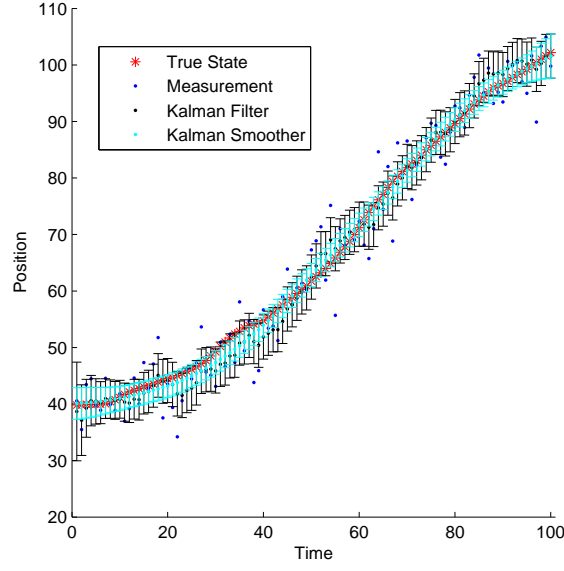


Figure 5: Kalman filter vs. Kalman smoother for the constant velocity model.

With prior mean  $\mu_0$  and covariance  $P_0$  we can write the full marginal log-likelihood as,

$$\log p(y) = \log N(y_1 \mid C\mu_0, R + CP_0C^T) + \sum_{t=1}^{T-1} \log N(y_{t+1} \mid CA\hat{x}_{t|t}, R + C(Q + AP_{t|t}A^T)C^T).$$

- c) See solution code in `em_lds.m`. Denote the Gaussian posterior distribution computed by the Kalman smoother by  $p(x_t \mid y) = N(x_t \mid \hat{x}_t, P_t)$ . From the course textbook or other resources, the  $M$ -step update of the covariance matrix  $R$  equals

$$R = \frac{1}{T} \sum_{t=1}^T (y_t - C\hat{x}_t)(y_t - C\hat{x}_t)^T + CP_tC^T.$$

The first term in the sum would be the ML estimate of  $R$  if  $x_t$  were observed to exactly equal its posterior mean  $\hat{x}_t$ . Because there is posterior uncertainty in the  $E$ -step, the covariance is increased by an amount proportional to the posterior covariance  $P_t$ .

- d) Figure 6(a) shows the tracking output for the **spiral** data based on parameters learned by EM. The log-likelihood bound versus EM iteration is in Figure 6(b), and the learned covariance is

$$R = \begin{bmatrix} 18.7408 & -1.5461 \\ -1.5461 & 21.0857 \end{bmatrix}.$$

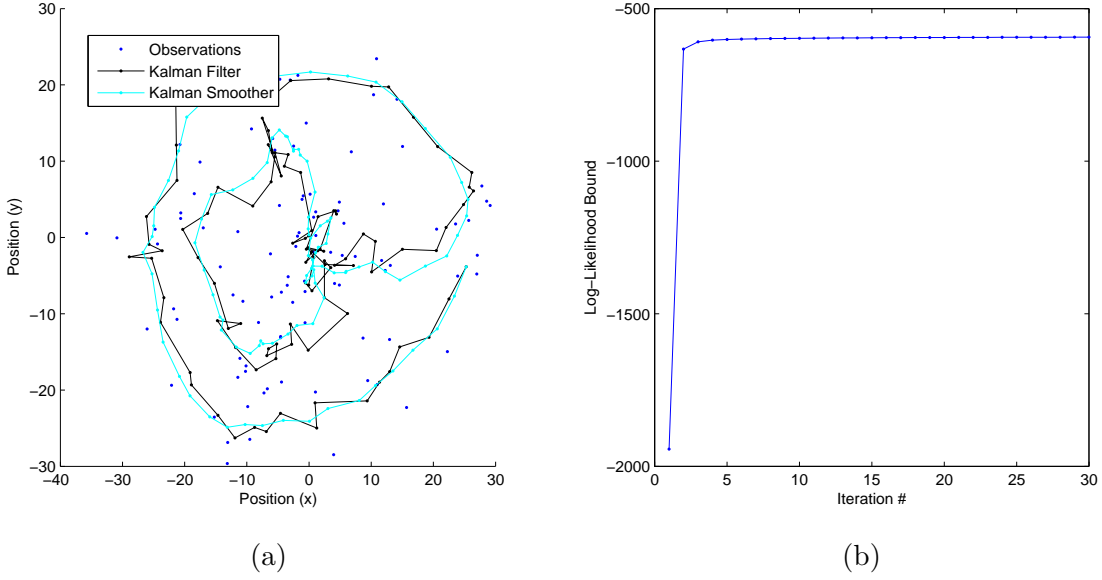


Figure 6: (a) Tracking output based on learned parameters for **spiral** data. (b) Log-likelihood bound vs. iteration. Note that it is monotonically non-decreasing.

### Question 3:

- a) Let  $N$  be the number of sequences in the training set and  $T_n$  the number of timescans in the  $n^{\text{th}}$  sequence. The transition parameters  $\pi \in \mathbb{R}^{K \times K}$  are computed by summing the total number of transitions from one state to another,

$$\tilde{\pi}_{ij} = \sum_{n=1}^N \sum_{t=2}^{T_n} \mathbb{I}(z_{n,t-1} = i) \mathbb{I}(z_{n,t} = j),$$

and then normalizing across columns,

$$\pi_{ij} = \frac{\tilde{\pi}_{ij}}{\sum_{k=1}^K \tilde{\pi}_{ik}}$$

Figure 7 shows the transition probabilities learned for each train/test split.

- b) Given state space model parameters the different training sequences are conditionally independent. We can then collect sufficient statistics for the E-step by running the Kalman smoother on each sequence with the current parameter estimates.

The M-step will depend on the sufficient statistics of all sequences. Where there is originally a summation over the length of a sequence  $\sum_{t=1}^T$  there will appear the double summation over sequences  $1, \dots, M$  and the corresponding length of that sequence  $T_m$  as in  $\sum_{m=1}^M \sum_{t=1}^{T_m}$ . For example, the dynamics matrix update would be,

$$A^{\text{new}} = \left( \sum_{m=1}^M \sum_{t=2}^{T_m} \mathbb{E}_m[x_t x_{t-1}^T] \right) \left( \sum_{m=1}^M \sum_{t=2}^{T_m} \mathbb{E}_m[x_{t-1} x_{t-1}^T] \right)^{-1},$$

where  $\mathbb{E}_m[\cdot]$  denotes the expected sufficient statistics of the  $m^{\text{th}}$  sequence.

	LT	RT	W
LT	0.97	0.00	0.027
RT	0.00	0.98	0.021
W	0.021	0.021	0.96

	LT	RT	W
LT	0.98	0.00	0.021
RT	0.00	0.98	0.020
W	0.014	0.013	0.97

	LT	RT	W
LT	0.98	0.00	0.025
RT	0.00	0.98	0.022
W	0.02	0.0172	0.96

Figure 7: Transition probabilities for each train/test split between “Left Turn” (LT), “Right Turn” (RT), and “Waggle” (W). Entry  $(i, j)$  in the matrix represents  $P(z_{n,t+1} = i \mid z_{n,t} = j)$ .

- c) *This is done in the code for you; it is the second output out of load\_bees.*
- d) *Figure 8 shows the log-likelihood bounds per iteration for each of the train/test splits. Learned parameter  $R$  for the first train/test split is,*

```

R_1 =
    0.0005    0.0002   -0.0027    0.0008
    0.0002    0.0004    0.0010    0.0001
   -0.0027    0.0010    0.0995   -0.0157
    0.0008    0.0001   -0.0157    0.0093
R_2 =
    0.0005   -0.0000    0.0008   -0.0006
   -0.0000    0.0006    0.0066   -0.0023
    0.0008    0.0066    0.1287   -0.0451
   -0.0006   -0.0023   -0.0451    0.0255
R_3 =
    0.0324   -0.0080   -0.0174   -0.0036
   -0.0080    0.0134   -0.0003    0.0034
   -0.0174   -0.0003    0.0139    0.0006
   -0.0036    0.0034    0.0006    0.0011

```




Figure 8: Log-likelihood bound per EM iteration on each of the three train/test splits of the dancing honeybees dataset.

e) The revised observation likelihood is given by,

$$p(\tilde{y} \mid x_t) = 0.9\text{Norm}(\tilde{y}_t \mid Cx_t, 0.1I_4 + R) + 0.1\text{Norm}(\tilde{y}_t \mid 0, 5I_4)$$

f) See solution code in `particle_filter_slds.m`.

g) Figures 9 and 10 show an analysis of the learning and inference on the dancing honeybee data for each of the three test sequences. In Figure 9 the ground truth (first row) 2D position and estimated 2D position (second row) show reasonable performance in predicting the target position on all test scenarios. On second and third test scenarios the inference underestimates the presence of “Waggle” state. The estimate angle is shown against ground truth (third row).

In Figure 10 the estimated state probabilities at each time step and the ground truth state are shown in the first row and second row, respectively. As can be seen, the ground truth discrete states are generally given high probability in the state sequence estimates.

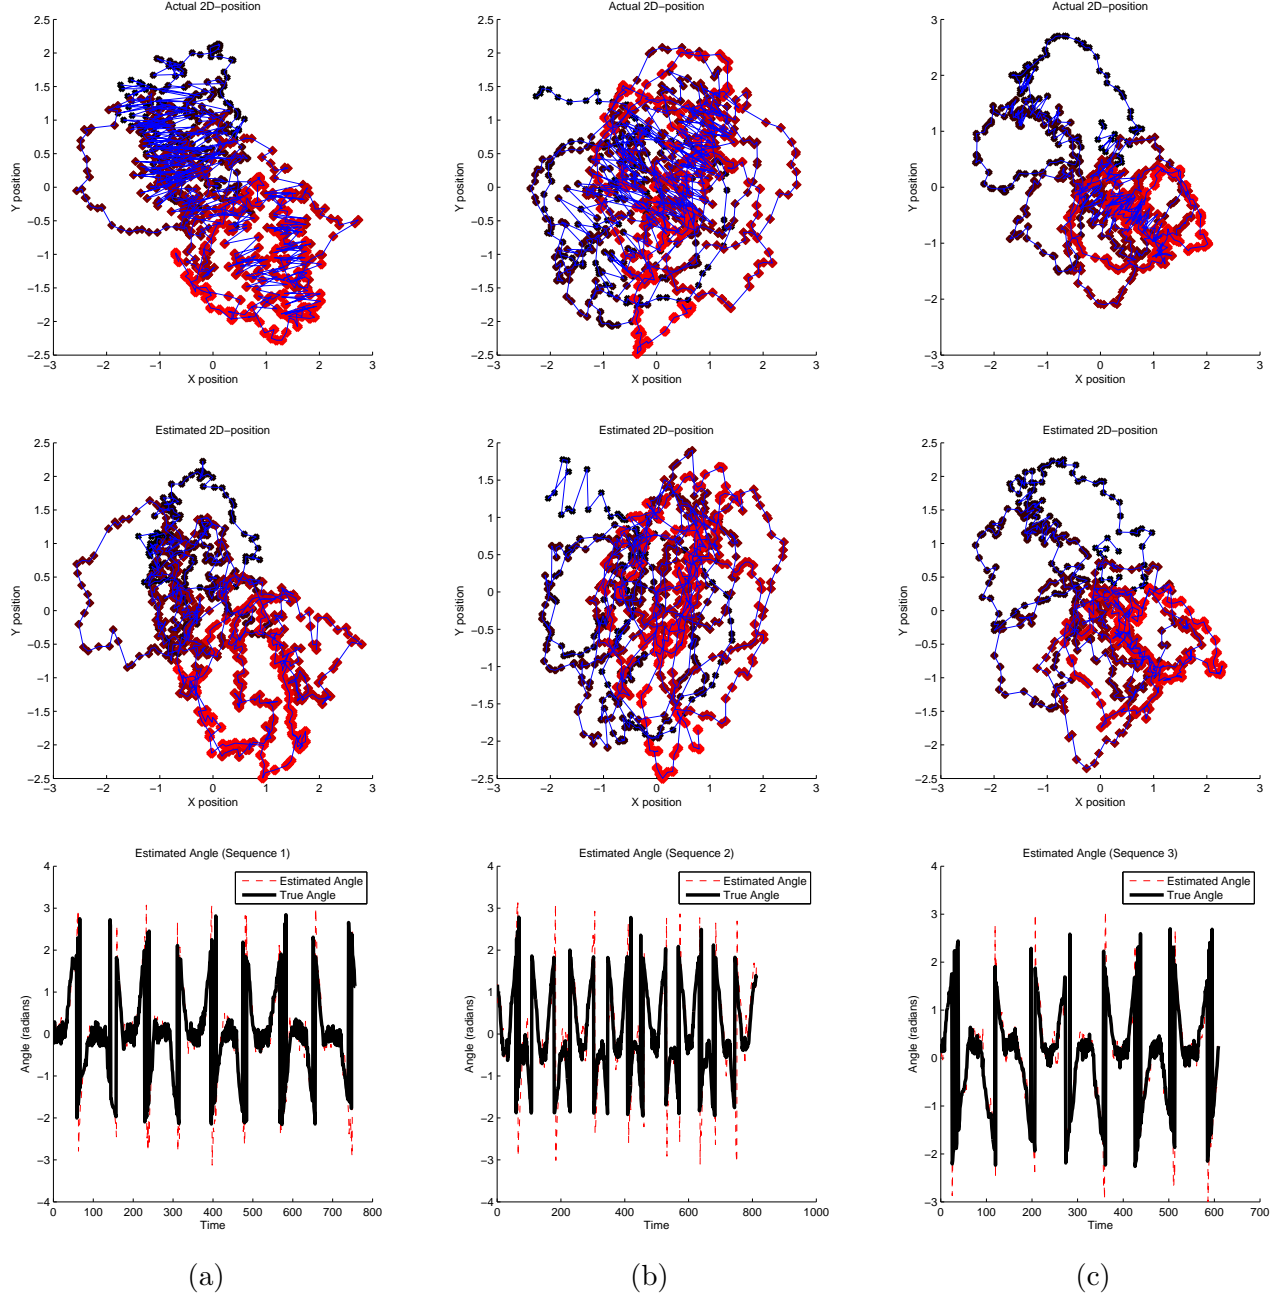


Figure 9: Analysis plots for the dancing honeybee data test sequences. (first row) True 2D position. (second row) Particle filter estimated position. (third row) Estimated angle vs. truth vs. time.



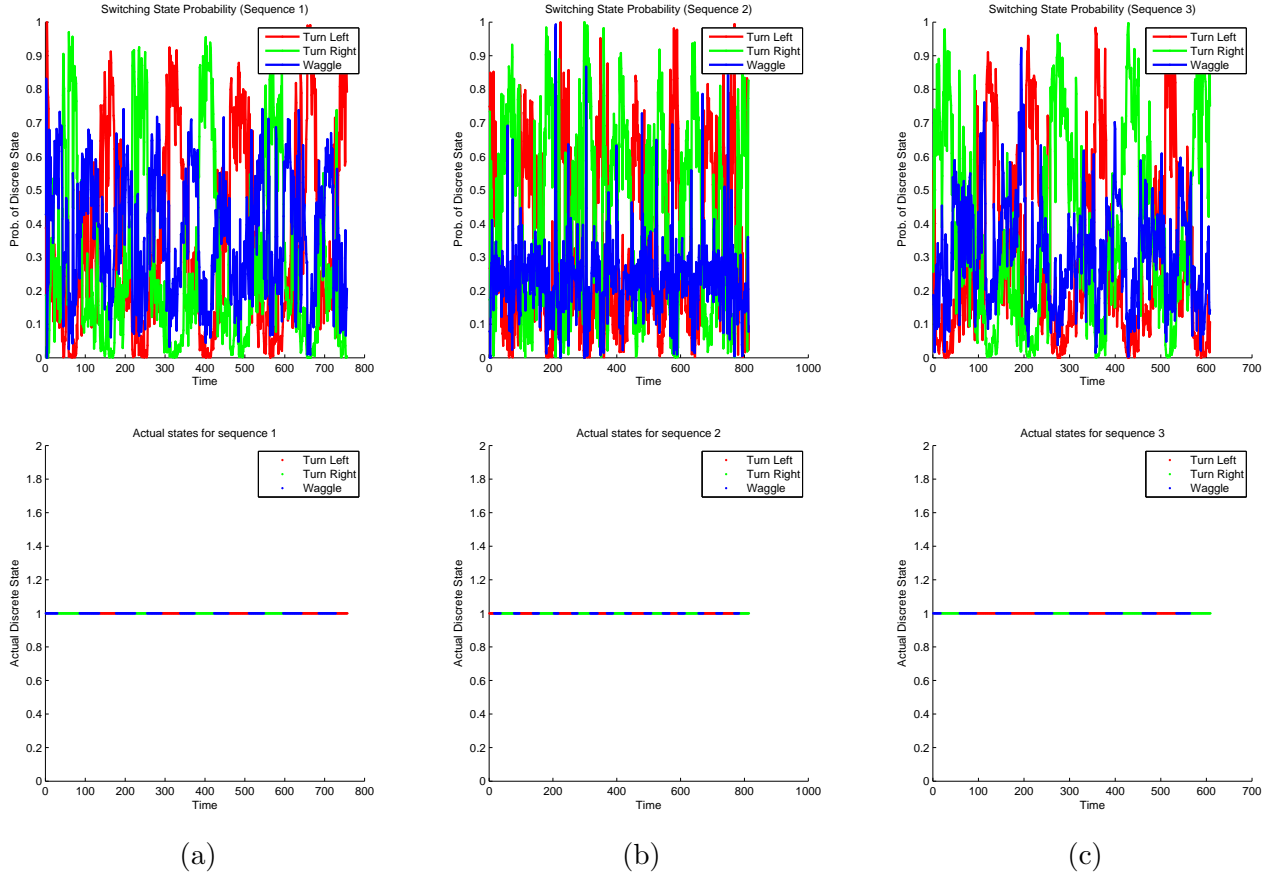


Figure 10: Analysis plots for the dancing honeybee data test sequences. Estimated probability of discrete state vs. time. (second row) Actual discrete state vs. time.